
HUMAN-ROBOT INTERACTION

ROBOT LEARNING

J. Verhaert

S1047220

Master Artificial Intelligence | Intelligent Technology
joost.verhaert@student.ru.nl

T. Gelton

S4480783

Master Artificial Intelligence | Intelligent Technology
thijs.gelton@student.ru.nl

1 Introduction

2 Methods and algorithms

The two algorithms that were used for the task of robot learning were a Feedforward Neural Network (FFN) and a Mixture Density Network (MDN). Regarding the popularity of deep learning, the FFN is widely used and does not need an indepth explanation. In short, it allows the modelling of non-linear distributions and performs exceptionally well on unstructured data. A MDN is a network that allows to make predictions under uncertainty. Meaning that a learned MDN outputs Gaussians and these are then used to sample from. This allows it to model many relations, where the input has multiple possible outputs. Bishop[1], shows an inverse sinusoid and how a MDN is able to model this where a FFN is not.

However, for any algorithm the first step to take is to gather data. For this we re-used some of Nao's parts from the first assignment, i.e. the keyboard functionality and the segmentation of the ball in the image. In a step-by-step manual fashion, the arm is first moved, then the ball is manually aligned and then 'E' is pressed to create a new data point. This CSV file is imported into a PyTorch Dataset object (see notebook) and deduplicated to make sure that every X and Y coordinate pair of the ball only has one corresponding roll and pitch angle.

After importing the data, which consisted of 72 samples after deduplication, the MDN was trained first. For this we used the Adam optimizer and a Negative Log-Likelihood (NLL) loss function. The Adam optimizer needed an epsilon of $1e-05$ to compensate for the often occurring vanishing gradient when using the NLL loss function. For this very reason, using a train and validation set was infeasible. The size of the dataset strongly impacted this vanishing gradient and thus all samples were needed in order to train the MDN model. Only one sample was reserved for testing. The network itself has two inputs that are connected, via 1 hidden layer of 6 nodes, to the π (π), σ (σ) and μ (μ) layer. Both the MDN and FFN use Tanh activation function. The number of Gaussians, and thus the number of π , σ and μ outputs, depends on the domain and it is up to the designer to specify this number. In the results (page 3) section we will compare the differences in the number of Gaussians.

To be able to compare the FFN and MDN, only one hidden layer with 6 nodes was used for the FFN as well. The network was then trained using the Mean Squared Error (MSE) and the same Adam optimizer as with the MDN.

3 Results

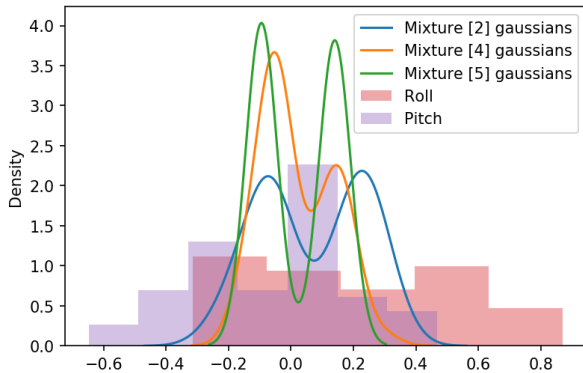


Figure 1: Mixture Models at different number of Gaussians and the target distribution. Input to each model was the mean of the x and y coordinates of the ball.

Eventually we chose to use 2, 4 and 5 gaussians and compare the results. The learned distributions can be seen in figure 1. In this specific case, where Nao's arm only used two degrees of freedom, the MDN would quickly converge to a point where the outputted π would become a one-hot encoded vector. This meant that the Mixture Model would converge to only one Gaussian. When inspecting figure 1, you can clearly see that there are two means for each mixture and that number of gaussians correlate to the size of the covariance.

Aside from inspecting Nao's behavior the FFN and the different MDN models were tested against the validation sample ($N = 1$). In figure 2 you can see that the best performing model is the MDN with 5 gaussians. Do note that this is not representative as it is only 1 sample, but in order to be able to train the MDNs without the vanishing gradient nearly all samples were needed for training. However, this one sample is unseen for all networks.

A factor that was impossible to capture by still images, was the actual performance of all models in the simulator. Eventually, the FFN seem to outperform the MDNs as the MDN caused some oscillation whereas the FFN tracked the ball in a seemingly perfect and smooth manner.

4 Conclusion

Learning is indispensable when it comes to intelligent systems.

References

- [1] C. M. Bishop, “Mixture density networks,” 1994.




Figure 2: Predictions of the MDN models and the FFN model on the validation sample ($N = 1$).